

Lab MA: Multi-Agent Systems

with LangChain & LangGraph

IST 488 — Building Human-Centered AI Applications

Time: ~40 minutes

Overview

In this lab, you will build a **Multi-Agent Trip Planner** as a Streamlit app. A **Supervisor** agent will orchestrate three specialist agents — *Research*, *Budget*, and *Itinerary* — to collaboratively plan a trip. No single agent can produce a complete plan; they must coordinate through the Supervisor.

By the end of this lab you will:

- Define custom tools using LangChain's `@tool` decorator
- Create specialist agents with `create_react_agent()`
- Wire them together under a Supervisor using `create_supervisor()`
- Build a Streamlit UI to interact with the multi-agent system
- Compare multi-agent output against a single-agent baseline
- Build a chatbot using `MessagesState` and `StateGraph`

Setup

1. Install the required packages:

```
pip install streamlit langchain langchain-openai langgraph
langgraph-supervisor
```

2. Make sure your OpenAI API key is configured in `.streamlit/secrets.toml`:

```
OPENAI_API_KEY = "sk-..."
```

3. You are provided a file called `travel_data.json` which contains all the destination information, cost data, and activity pools. Place this file in your `LABS-IST488` directory.
4. Create a new file called `Lab_MA.py` in your `LABS-IST488` directory.
5. Add `Lab_MA.py` to your `streamlit_app.py` navigation so you can access it.

1 Part 1: Setup & Environment (~5 minutes)

At the top of your `Lab_MA.py` file, add the following imports:

```
import os
import json
import streamlit as st
from langchain_openai import ChatOpenAI
from langchain_core.tools import tool
from langgraph.prebuilt import create_react_agent
from langgraph_supervisor import create_supervisor
```

Then set up the page:

- Add a title: “Lab MA: Multi-Agent Trip Planner”
- Add a brief description explaining how the system works (a Supervisor coordinates three specialist agents: Research, Budget, Itinerary)
- Load the OpenAI API key from `st.secrets` (same pattern as earlier labs)
- Create **two** `ChatOpenAI` instances:

```
agent_llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)
supervisor_llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)
```

When complete, run `streamlit run streamlit_app.py` and verify your page loads with the title and description.

2 Part 2: Building Specialist Agents (~10 minutes)

In this section you will define three tools and three agents.

Step 2A: Load the Travel Data

Load the `travel_data.json` file at the top of your tools section:

```
with open("travel_data.json", "r") as f:
    TRAVEL_DATA = json.load(f)
```

The JSON file has this structure:

- **destinations:** dict of city names → {highlights, best_time, culture, tips, weather}
- **flight_estimates:** dict of city names → estimated roundtrip flight cost (USD)
- **daily_costs:** dict of budget levels → {accommodation, food, transport, activities, misc}
- **activities:** dict of interest categories → list of activity strings
- **money_saving_tips:** list of tip strings

Step 2B: Define the Tools

Each tool is a Python function decorated with `@tool`. The decorator automatically registers the function’s name, docstring, and parameter types so that an LLM agent can call it:

```
@tool
def my_tool_name(param1: str, param2: int) -> str:
```

```

"""Description of what this tool does. The LLM reads this."""
# ... your logic ...
return json.dumps(result)

```

Create **three** tools that pull data from TRAVEL_DATA:

1. `search_destination(query: str) -> str`
Purpose: Look up travel info about a destination.
Implementation: Read `TRAVEL_DATA["destinations"]`. Match the user's query against city names (case-insensitive). Return the matching city's data as a JSON string. If no city matches, return generic fallback info.
2. `calculate_budget(destination: str, days: int, budget_level: str) -> str`
Purpose: Estimate trip costs.
Implementation: Read `TRAVEL_DATA["daily_costs"]` and `TRAVEL_DATA["flight_estimates"]`. Look up the daily rate for the given budget level, calculate totals based on the number of days, add the flight cost. Include `TRAVEL_DATA["money_saving_tips"]`. Return as JSON.
3. `create_schedule(destination: str, days: int, interests: str) -> str`
Purpose: Build a day-by-day itinerary.
Implementation: Read `TRAVEL_DATA["activities"]`. Parse the interests string, collect matching activities from the pool, and build a schedule with morning/afternoon/evening slots for each day. Return as JSON.

TIP

These tools return **simulated data** from the JSON file — you are not calling real APIs. This keeps the lab self-contained and predictable. The LLM still reasons about the data and composes a natural-language response.

Step 2C: Create the Agents

Use `create_react_agent()` to create three specialist agents:

```

agent = create_react_agent(
    model=agent_llm,
    tools=[your_tool],
    name="agent_name",
    prompt="You are a specialist in ... Always use the tool ..."
)

```

Create these three agents:

1. **research_agent** — Tools: `[search_destination]`
 Prompt: Tell it to be a travel research specialist that always uses the `search_destination` tool.
2. **budget_agent** — Tools: `[calculate_budget]`
 Prompt: Tell it to be a budget specialist that always uses the `calculate_budget` tool.
3. **itinerary_agent** — Tools: `[create_schedule]`
 Prompt: Tell it to be an itinerary specialist that always uses the `create_schedule` tool.

IMPORTANT

In each agent's prompt, instruct it to **ALWAYS** call its tool and not make up information. This ensures the agent uses the tool rather than relying on its own knowledge.

3 Part 3: Creating the Supervisor (~8 minutes)

The Supervisor is the central orchestrator. It receives the user's request, decides which agent(s) to delegate to, and synthesizes the final response.

```
workflow = create_supervisor(
    agents=[research_agent, budget_agent, itinerary_agent],
    model=supervisor_llm,
    prompt="..."
)
```

Write a supervisor prompt that:

- Lists the three available agents and what each one does
- Instructs the supervisor to route questions:
 - Destination questions → `research_agent`
 - Cost/budget questions → `budget_agent`
 - Schedule/itinerary questions → `itinerary_agent`
 - Full trip planning → **ALL** three agents
- Tells the supervisor to synthesize all responses into one organized plan

Then compile and invoke:

```
multi_agent_app = workflow.compile()

result = multi_agent_app.invoke({
    "messages": [{"role": "user", "content": "your query here"}]
})

# The final response:
result["messages"][-1].content
```

4 Part 4: Streamlit Interface (~10 minutes)

Build the user interface to interact with your multi-agent system.

Step 4A: Sidebar Controls

In a `with st.sidebar:` block, create:

- A `text_input` for the destination (default: "Paris")
- A `slider` for trip duration (1–14 days, default: 5)
- A `selectbox` for budget level (Budget / Moderate / Luxury)
- A `multiselect` for interests (Food, History, Art, Nature, etc.)

Step 4B: Session State

Initialize session state variables to persist results across reruns:

```
if "ma_result" not in st.session_state:
    st.session_state.ma_result = None
if "ma_messages" not in st.session_state:
    st.session_state.ma_messages = None
```

Step 4C: Query Construction & Invocation

Build a query string from the sidebar inputs, e.g.:

“Plan a 5-day trip to Paris. My budget level is moderate. My interests include: Food, History. Please provide destination research, a budget breakdown, and a day-by-day itinerary.”

Create a “Plan My Trip” button. When clicked: show a spinner, invoke `multi_agent_app`, store the result in session state, and display it with `st.markdown()`.

Step 4D: Agent Activity Log

After displaying results, add a debug section in the sidebar showing which agents were called and in what order:

```
for msg in result["messages"]:
    msg_name = getattr(msg, "name", None) # agent name
    tool_calls = getattr(msg, "tool_calls", None) # tool invocations
```

Display the execution order with emoji labels (e.g., `research_agent`, `budget_agent`, `itinerary_agent`).

5 Part 5: Single-Agent Comparison (~7 minutes)

Add a section below the trip plan that compares multi-agent output with a single-agent response.

1. Add a divider and subheader: “Experiment: Single Agent vs Multi-Agent”
2. Add a button “Run Single-Agent Comparison”. When clicked, send the **exact same query** to a single LLM with no tools:

```
single_llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)
single_result = single_llm.invoke(trip_query)
st.markdown(single_result.content)
```

3. Compare the two responses side-by-side. Note the differences in data quality, structure, and specificity. You will reflect on these observations in the Write-Up Questions at the end of this document.

6 Part 6: Chatbot Mode — MessagesState & StateGraph (~10 min bonus)

In Parts 1–5 you used `create_supervisor()`, a **high-level** helper. Now you will build a conversational chatbot using the **lower-level** LangGraph primitives.

Key Concepts

MessagesState

A `TypedDict` that holds a list of messages as the graph’s shared state. Every node reads from and writes to this state. Think of it as the “conversation memory” flowing through the graph.

```
from langgraph.graph import MessagesState
# Equivalent to:
# class MessagesState(TypedDict):
#     messages: list[BaseMessage]
```

StateGraph

The core graph builder. You define **nodes** (processing functions) and **edges** (connections), then compile into a runnable app.

```
from langgraph.graph import StateGraph, START, END
```

Nodes

Python functions that take the current state and return updated messages:

```
def my_node(state: MessagesState):
    messages = state["messages"] # read
    # ... process ...
    return {"messages": [response]} # write
```

Edges

Three types of connections:

- **Static:** `graph.add_edge("node_a", "node_b")`
- **Conditional:** `graph.add_conditional_edges("node_a", router_fn, mapping)`
- **START/END:** `graph.add_edge(START, "first_node")`

Step 6A: Define Graph Nodes

Create these node functions:

1. **supervisor_node** — reads messages, uses the LLM with a routing prompt to output one word: “research”, “budget”, “itinerary”, “all”, or “chat”
2. **research_node** — invokes the `research_agent` from Part 2
3. **budget_node** — invokes the `budget_agent`
4. **itinerary_node** — invokes the `itinerary_agent`
5. **run_all_agents** — sequentially invokes all three agents, combines outputs
6. **synthesizer_node** — produces a final, well-organized response from all agent outputs

Step 6B: Define the Router

```
def route_from_supervisor(state: MessagesState):
    last_msg = state["messages"][-1].content.strip().lower()
    if "research" in last_msg:
        return "research"
    elif "budget" in last_msg:
        return "budget"
    elif "itinerary" in last_msg:
        return "itinerary"
    elif "all" in last_msg:
        return "all"
    else:
        return "chat"
```

Step 6C: Build the StateGraph

```
chatbot_graph = StateGraph(MessagesState)

# Add nodes
chatbot_graph.add_node("supervisor", supervisor_node)
chatbot_graph.add_node("research", research_node)
chatbot_graph.add_node("budget", budget_node)
chatbot_graph.add_node("itinerary", itinerary_node)
chatbot_graph.add_node("all_agents", run_all_agents)
chatbot_graph.add_node("synthesizer", synthesizer_node)

# Edges
chatbot_graph.add_edge(START, "supervisor")

chatbot_graph.add_conditional_edges(
    "supervisor",
    route_from_supervisor,
    {
```

```

        "research": "research",
        "budget": "budget",
        "itinerary": "itinerary",
        "all": "all_agents",
        "chat": "synthesizer",
    },
)

chatbot_graph.add_edge("research", "synthesizer")
chatbot_graph.add_edge("budget", "synthesizer")
chatbot_graph.add_edge("itinerary", "synthesizer")
chatbot_graph.add_edge("all_agents", "synthesizer")
chatbot_graph.add_edge("synthesizer", END)

chatbot_app = chatbot_graph.compile()

```

The resulting graph:

```

START
|
supervisor
|
+----> research ----+
+----> budget -----+----> synthesizer ---> END
+----> itinerary ----+
+----> all_agents ---+
+----> synthesizer --+

```

Step 6D: Streamlit Chat Interface

Build a chat interface using `st.chat_input` and `st.chat_message`:

1. Initialize session state for chat history
2. Display existing chat history (loop through messages)
3. On new user input: append to state, invoke `chatbot_app`, extract and display response
4. Try these flows:
 - “Tell me about Tokyo” (routes to research)
 - “How much would that cost for 5 days?” (routes to budget)
 - “Plan me a full trip to Paris” (routes to all agents)
 - “Thanks!” (routes to chat/general)

Testing Your App

Run your completed app: `streamlit run streamlit_app.py`

Verify the following:

- ☐ The page loads with the title and description
- ☐ Sidebar controls work (destination, days, budget, interests)
- ☐ Clicking “Plan My Trip” produces a comprehensive trip plan
- ☐ The Agent Activity Log shows all three agents were called
- ☐ The single-agent comparison produces a noticeably different response

- ☐ Changing the destination and interests changes the output
- ☐ The Part 6 chatbot responds conversationally
- ☐ The chatbot correctly routes to different agents based on the question
- ☐ The chatbot remembers previous messages in the conversation

Write-Up Questions

Answer the following questions in a separate document (submitted via Blackboard or as a text file). Each response should be 3–5 sentences unless otherwise specified.

Question 1: Agent Architecture

Explain the Supervisor (orchestrator/subagent) pattern used in this lab. What role does the Supervisor play, and how does it differ from the specialist agents? Why is this pattern useful for complex tasks?

Question 2: Single-Agent vs. Multi-Agent Comparison

Compare the multi-agent trip plan to the single-agent response.

- Which response had more specific, actionable information? Why?
- Which response was better organized? What caused the difference?
- Identify at least TWO concrete examples where the multi-agent response contained data that the single-agent response lacked.

Question 3: Trade-offs

Multi-agent systems are not always better. Discuss at least THREE trade-offs or downsides of using a multi-agent approach compared to a single-agent approach. Consider: latency, cost, complexity, and failure modes.

Question 4: MessagesState & StateGraph

In Part 6 you built a chatbot using `StateGraph` and `MessagesState`.

- What is `MessagesState` and what role does it play in the graph?
- Explain the difference between a static edge (`add_edge`) and a conditional edge (`add_conditional_edges`). Give an example of each from your code.
- How does `create_supervisor()` from Parts 1–5 relate to the `StateGraph` you built manually? What does it do for you automatically?

Submission Checklist

- ☐ `Lab_MA.py` — runs without errors, includes Parts 1–6
- ☐ Write-up with answers to all 4 questions